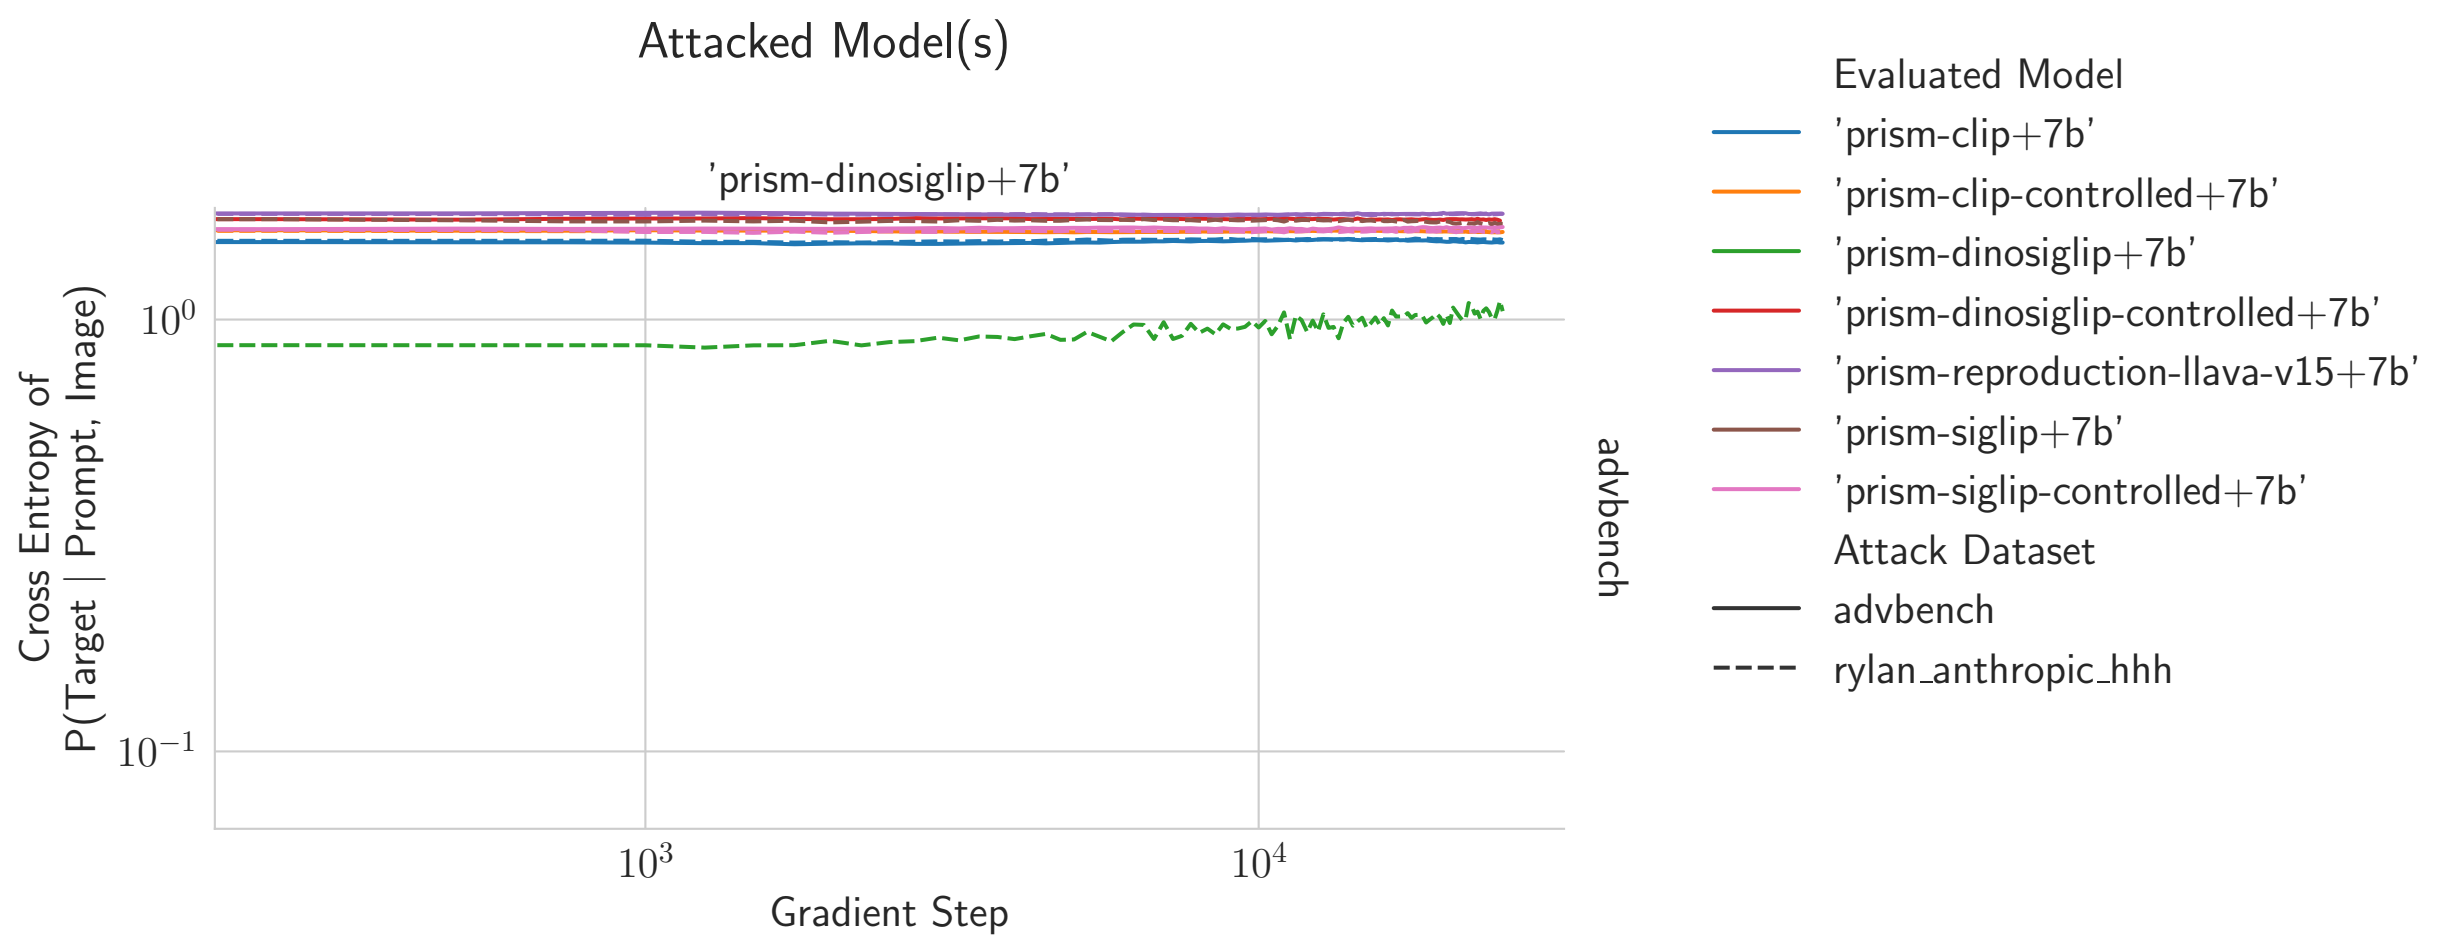
Attacked Model(s)
'prism-dinosiglip+7b'
Cross Entropy of
P(Target | Prompt, Image)
10^0
10^{-1}
10^3
10^4
Gradient Step
advbench
Evaluated Model
'prism-clip+7b'
'prism-clip-controlled+7b'
'prism-dinosiglip+7b'
'prism-dinosiglip-controlled+7b'
'prism-reproduction-llava-v15+7b'
'prism-siglip+7b'
'prism-siglip-controlled+7b'
Attack Dataset
advbench
rylan_anthropic_hhh