

Attacked Model(s)



advbench

Evaluated Model

- 'prism-clip+7b'
- 'prism-clip-controlled+7b'
- 'prism-reproduction-llava-v15+7b'
- 'prism-siglip+7b'
- 'prism-siglip-controlled+7b'

Attack Dataset

- advbench
- rylan_anthropic_hhh