

Attacked Model(s)

'prism-clip-controlled+7b'

Cross Entropy of
P(Target | Prompt, Image)

0.0

0.5

1.0

1.5

0

5000

10000

15000

20000

25000

Gradient Step

advbench

Evaluated Model

'prism-clip+13b'

'prism-clip+7b'

'prism-clip-controlled+7b'

'prism-dinosiglip+7b'

'prism-dinosiglip-controlled+7b'

'prism-reproduction-llava-v15+13b'

'prism-reproduction-llava-v15+7b'

'prism-siglip+7b'

'prism-siglip-controlled+13b'

'prism-siglip-controlled+7b'

Attack Dataset

advbench

rylan_anthropic_hhh