


Attacked Model(s)

'prism-siglip+7b'



Evaluated Model

- 'prism-clip+7b'
 - 'prism-clip-controlled+7b'
 - 'prism-dinosiglip+7b'
 - 'prism-dinosiglip-controlled+7b'
 - 'prism-reproduction-llava-v15+7b'
 - 'prism-siglip+7b'
 - 'prism-siglip-controlled+7b'
- ## Attack Dataset
- advbench
 - rylan_anthropic_hhh